


DeepSeekCell evidence-guided selective refinement

One cached first-pass LLM output is audited by deterministic biological evidence before any second-pass inference is spent.



Paired ablation benchmark

All DeepSeekCell arms reuse the same cached first-pass response hash. Plain, Evidence and Calibrated keep labels fixed; selector arms spend matched second-pass budgets.

Random-k

Confidence-k

NoOntology-k

Evidence-k

FullRefined